

Default Project 1 - DD2424, 2023

Teachers and TAs of DD2424 2023 and earlier

1 Explore Transfer Learning

The first default project will explore the concept of transfer learning. This is one of the most common use cases of deep learning - download a pre-trained model and then adapt it to your dataset. Please visit the tutorial [FINETUNING TORCHVISION MODELS](#), this gives an overview of how to perform fine-tuning of with a pre-trained ConvNet within PyTorch. Another tutorial is given at [FastAI Computer Vision tutorial](#).

1.1 Basic project to get E

- Download a pre-trained modern ConvNet such as: ResNet18, ResNet34,...
- Download the dataset [The Oxford-IIIT Pet Dataset](#)
- Replace the final layer of the pre-trained ConvNet to solve the binary classification problem of recognising pictures of Dog Vs Cat. Fine-tune the replaced final layer with the Pet Dataset's training data. Use Adam or NAG optimizer. Without too much effort you should be able to get very high performance ($\geq 99\%$ test accuracy) on this binary classification task.

Note you will have to check what spatial size of input your pre-trained network can cope with. The default ResNet architectures have a global average pooling layer, just before the final fully connected output layer, that produces a feature vector of fixed size independent of the spatial extent of the input image. Thus you just have to ensure that the amount of down-sampling implemented during the ResNet does not make the spatial extent of the feature maps disappear for the later layers given the size of your input images. If your images are too small and this will happen then you should re-size them to the smallest acceptable size so that your computational effort is minimized. The normal procedure is to resize the image with one scale factor, to maintain the image's aspect ratio, so its shortest side is re-scaled to the target length.

- Next your goal is to solve the multi-class classification problem of recognising the breed of cat or dog. In this case you have to replace the final layer to have 37 outputs. As this multi-class problem is harder than the binary classification problem you have just solved, you will have to do more work and you should fine-tune more of the network as opposed to just the replaced final layer. You should explore each of the following issues when fine-tuning:
 - Benefit of fine-tuning more layers (from just the final layer, and then incrementally increasing the number of layers fine-tuned).
 - Different learning rates and/or learning rate schedulers for different layers

- Benefit of applying data augmentation during training (flip, small rotations, crops, small size scaling)
- Effect of fine-tuning or not the batch-norm parameters and updating the estimate of the batch mean and standard deviations on the final performance on the new dataset.

After these experiments you should be able to get a final test accuracy of $\geq 95\%$. You do not have to use the Pet Dataset. You are free to use another dataset but it should have comparable or greater difficulty than the Pet Dataset w.r.t. number of classes and size of images.

1.2 Extending the basic project to get a higher grade

Once you have explored relatively thoroughly the basic project then your group can add extensions to aim for a higher grade.

1.2.1 From E \rightarrow D/C

If you are aiming for a D or C then here are some extensions you could apply to investigate if it was possible to improve performance.

- Explore using deeper networks than you used in the basic project. Does the deeper network help? Is it trickier to fine-tune? Do you even need to fine-tune the earlier layers? Do you need to change the optimizer to for instance AdamW to ensure good training?
- Change the loss you use for the multi-classification problem to say multi-class BCE training or multi-class SVM etc. Does your new loss improve the final test accuracy or speed up convergence etc?
- Just fine-tune the batch norm mean and standard deviation and keep the weights of the layers (except the final layer) frozen and see if you can improve results.
- Add more sophisticated data-augmentations such as random erasing, CutMix or MixUp. to help with regularization.

Assuming the write-up and presentation of the material is good, how many extensions are needed to achieve the desired grade? This is a tricky question as exploring one extension in depth can be considered equivalent to exploring many extensions superficially. So for example exploring CutMix and its benefits thoroughly should be more than sufficient to get a C, while a cursory investigation of the benefits of a deeper network and multi-class BCE training may just about make the cut for a D grade.

You are, of course, encouraged to come up with your own extensions. But do remember to get them vetted through your project proposal.

1.2.2 From E \rightarrow B/A

If you are aiming for a B or A then here are some possible extensions from the basic project you could apply to investigate if it was possible to improve performance.

- **Coping with fine-tuning with fewer labelled training examples** Try decreasing the percentage of the labelled training data used during fine-tuning - all the training data, 50%, 10%, 1%. Keep a record of the drop in performance as the percentage of training data used is dropped. In this extension you can explore whether it is possible to use only a small percentage of the training data with their labels plus the rest of the training data without their labels which should be considered as the unlabelled training data set and get the same performance as when you use the full labelled training dataset during fine-tuning. Using both labelled and unlabelled training datasets is known as semi-supervised learning and there has been an explosion of methods and approaches to learn under this regime. The following blog gives a very extensive and readable overview: [Learning with not Enough Data Part 1: Semi-Supervised Learning](#). For your project you are, of course, only expected to explore one approach or paper such as one focusing on
 - Some form of consistency regularization or
 - Pseudo-labelling
 - etc...
- There are pre-trained Vision Transformers (ViTs) available for example at Hugging Face [Vision Transformer \(ViT\)](#) and descriptions of how to fine-tune at this blog [Fine-Tune ViT for Image Classification with HuggingFace Transformers](#). You can build on this example with the basic project in mind. The report could also highlight the advantages/disadvantages and the level of difficulty of fine-tuning a ResNet.

You are, of course, encouraged to come up with your own extensions. But do remember to get them vetted through your project proposal. If you do go for an extension from an E to an A then most of the project report should be devoted to the extension as opposed to the basic project. For the basic assignment you should report the main results and put the more extensive fine-tuning results in the appendix.